

IA Paris 8 - 9 février 2026

Simon Willison's Weblog

[Subscribe](#)

[theatre](#) 4 [ai](#) 1913
[generative-ai](#) 1605 [llms](#) 1971
[ai-assisted-programming](#) 3160
[tiktok](#) 12 [ai-ethics](#) 258
[coding-agents](#) 137 [claude-code](#) 82

Monthly briefing

Sponsor me for \$10/month and get a curated email digest of the month's most important LLM developments.

Pay me to send you less!

[Sponsor & subscribe](#)

We gotta talk about AI as a programming tool for the arts. Chris Ashworth is the creator and CEO of [QLab](#), a macOS software package for "cue-based, multimedia playback" which is designed to automate lighting and audio for live theater productions.

I recently started following him on TikTok where he posts about his business and theater automation in general - Chris founded [the Voxel theater](#) in Baltimore which QLab use as a combined performance venue, teaching hub and research lab (here's [a profile of the theater](#)), and the resulting videos offer a fascinating glimpse into a world I know virtually nothing about.

[This latest TikTok](#) describes his Claude Opus moment, after he used Claude Code to build a custom lighting design application for a very niche project and put together a useful application in just a few days that he would never have been able to spare the time for otherwise.

Chris works full time in the arts and comes at generative AI from a position of rational distrust. It's interesting to see him working through that tension to acknowledge that there are valuable applications here to build tools for the community he serves.

I have been at least gently skeptical about all this stuff for the last two years. Every time I checked in on it, I thought it was garbage, wasn't interested in it, wasn't useful. [...] But as a programmer, if you hear something like, this is changing programming, it's important to go check it out once in a while. So I went and checked it out a few weeks ago. And it's different. It's astonishing. [...]

One thing I learned in this exercise is that it can't make you a fundamentally better programmer than you already are. It can take a person who is a bad programmer and make them faster at making bad programs. And I think it can take a person who is a good programmer and, from what I've tested so far, make them faster at making good programs. [...] You see programmers out there saying, "I'm shipping code I haven't looked at and don't understand." I'm terrified by that. I think that's awful. But if you're capable of understanding the code that it's writing, and directing, designing, editing, deleting, being quality control on it, it's kind of astonishing. [...]

The positive thing I see here, and I think is worth coming to terms with, is this is an application that I would never have had time to write as a professional programmer. Because the audience is three people. [...] There's no way it was worth it to me to spend my energy of 20 years designing and implementing software for artists to build an app for three people that is this level of polish. And it took me a few days. [...]

I know there are a lot of people who really hate this technology, and in

<https://simonwillison.net/2026/Jan/30/a-programming-tool-for-the-arts/>

2026-02-03 16:02:12

Home [Blog](#)

[GitHub](#) [Twitter](#) 

A field guide to sandboxes for AI

Jan 5, 2026



Photo by [Christian Weiss](#)

Every AI agent eventually asks for the same thing:

“
Let me run a program.

<https://www.luiscardoso.dev/blog/sandboxes-for-ai>

2026-02-03 16:11:28

embedding-shapes

Home Posts About

One Human + One Agent = One Browser From Scratch

2026-01-27

Just for the fun of it, I thought I'd embark on a week-long quest to generate millions of tokens and millions of lines of source code to create one basic browser that can render HTML and CSS (no JS tho), and hopefully I could use this to receive even more VC investments.

But then I remembered that I have something even better: a human brain! It is usually better than any machine at coordinating and thinking through things, so let's see if we can hack something together, one human brain and one LLM agent brain!



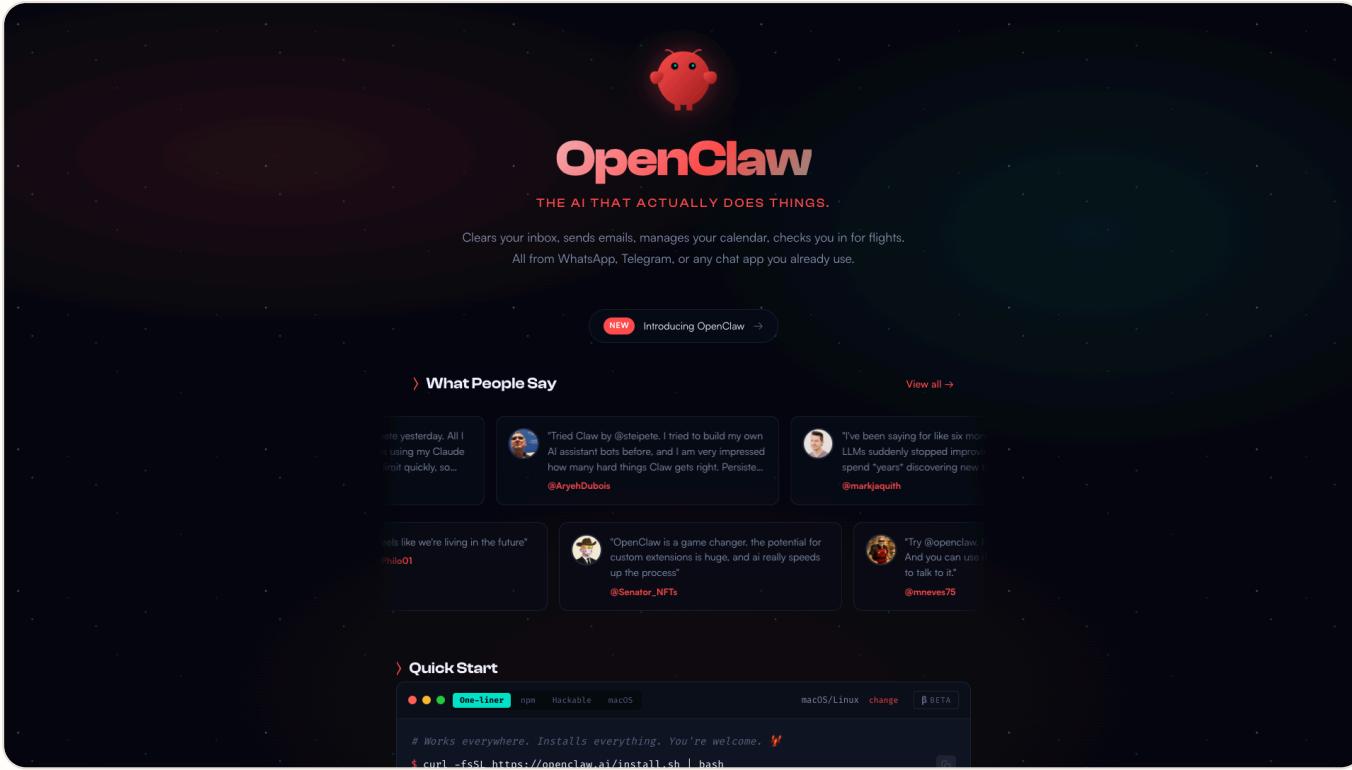
The above might look like a simple .webm video, but it's actually a highly sophisticated and advanced browser that was super hard to build, encoded as pixels in a video file! Wowzers.

Day 1 - Starting out

For extra fun when building this, I set these requirements for myself and the agent:

<https://emsh.cat/one-human-one-agent-one-browser/>

2026-02-03 16:11:35



<https://openclaw.ai/>

2026-02-03 16:11:41

Running OpenClaw in Docker

I'm not brave enough to run [OpenClaw](#) (aka Clawbot aka Moltbot) directly on my Mac, so I decided to try running it in a Docker instead container.

OpenClaw has Docker support out of the box, described [on this page](#) of their documentation. Here's what worked for me.

Use their Docker Compose configuration

First, clone their GitHub repository:

```
git clone https://github.com/openclaw/openclaw
```

It contains a script for running OpenClaw in Docker called [docker-setup.sh](#) which in turn uses Docker Compose and their [docker-compose.yml](#) file.

The script will create two folders directly on your Mac which will then be mounted as volumes in the Docker container:

- `~/openclaw` is the configuration directory. This will eventually contain OpenClaw memory, configuration, third party API keys, etc.
- `~/openclaw/workspace` is the workspace directory full of files that are directly available to the agent as it runs inside the container. Files the agent creates will be saved here too.

Answering all of those questions

On first run OpenClaw asks you a lot of questions. Most of these are reasonably obvious but I still had to start over a couple of times to get everything right. Some that I found tricky:

- Onboarding mode: manual
- What do you want to set up?: Local gateway (this machine)
- Model provider. I decided to go for OpenAI Codex with ChatGPT OAuth, which then allowed me to authenticate against ChatGPT to spend tokens already covered by my \$20/month OpenAI subscription. I did this because I've heard that OpenClaw can spend a lot of tokens on API plans, and using ChatGPT put an easy upper limit on how much it could spend. When you opt for this OpenClaw gives you a URL to open in your browser which redirects back to a non-running `localhost` service and displays an error message - you then copy and paste that `localhost` URL back into OpenClaw to complete the authentication. Here's [what that looks like](#).
- Tailscale: I tried to configure this the first time and it resulted in a machine I couldn't use, so the second time I said "no".

Once it's up and running you can run:

```
docker ps
```

To see the container. Mine is running the image `openclaw:local` and has a container name `openclaw-openclaw-1`

<https://til.simonwillison.net/llms/openclaw-docker>

2026-02-03 16:11:48

DANIEL MIESSLER

HOME BLOG TELOS IDEAS PROJECTS PREDICTIONS ABOUT MEMBERS UL SITE DAEMON Q ...

Building a Personal AI Infrastructure (PAI)

How I built my own unified, modular Agentic AI system named Kai

July 26, 2025

#ai #technology #future #innovation

Updated — January 2026

This post has been completely rewritten for PAI v2.4 (January 2026). The architecture framework, Algorithm, Memory System, and Hook System are all new. The [ORIGINAL VERSION](#) of this post (July–December 2025) is preserved separately.

Pioneering PAI: How Daniel Miessler's Personal AI Infrastructure... Partager

BUILDING A PERSONAL AI SYSTEM

Regarder sur YouTube

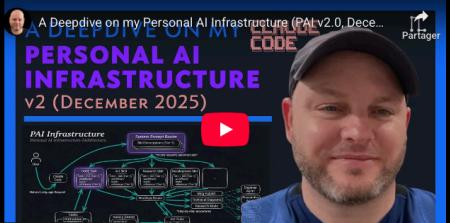
DANIEL MIESSLER ON THE COGNITIVE REVOLUTION DISCUSSING PAI AND THE FUTURE OF PERSONAL AI

A Deepdive on my Personal AI Infrastructure (PAI) v2.0, Dec... Partager

PERSONAL AI INFRASTRUCTURE

v2 (DECEMBER 2025)

PAI Infrastructure



l'IA n'est pas "magique", c'est une architecture socio-technique (prompts, outils, sandbox, permissions, coûts)

<https://danielmiessler.com/blog/personal-ai-infrastructure>

2026-02-03 16:16:09

Peter Steinberger

Posts About Q ⌂

Shipping at Inference-Speed

28 Dec, 2025 • 18 min read | Edit on GitHub

What Changed Since May

It's incredible how far "vibe coding" has come this year. Whereas in May I was amazed that *some* prompts produced code that worked out of the box, *this is now my expectation*. I can ship code now at a speed that seems unreal. I burned a lot of tokens since then. Time for an update.

It's funny how these agents work. There's been this argument a few weeks ago that *one needs to write code in order to feel bad architecture* and that using agents creates a disconnection – and I *couldn't disagree more*. When you spend enough time with agents, you know exactly how long sth should take, and when codex comes back and hasn't solved it in one shot, I already get suspicious.

The amount of software I can create is now mostly *limited by inference time and hard thinking*. And let's be honest – most software does not require hard thinking. Most apps shove data from one form to another, maybe store it somewhere, and then show it to the user in some way or another. The simplest form is text, so by default, whatever I wanna build, it starts as CLI. Agents can call it directly and verify output – closing the loop.

The Model Shift

The real unlock into building like a factory was GPT 5. It took me a few weeks after the release to see it – and for codex to catch up on features that claudie code had, and a bit to learn and understand the differences, but then I started trusting the model more and more. *These days I don't read much code anymore*. I watch the stream and sometimes look at key parts, but I gotta be honest – most code I don't read. I do know where which components are and how things are structured and how the overall system is designed, and that's usually all that's needed.

<https://steipete.me/posts/2025/shipping-at-inference-speed>

2026-02-03 16:16:26

A screenshot of a blog post titled "Singing the gospel of collective efficacy" by Matt Webb. The post discusses collective efficacy in a neighborhood, mentioning swift nest boxes and Greta Thunberg's sense of collective efficacy.

A BLOG BY MATT WEBB
[About Archive Work](#)

SUBSCRIBE FOR \$0
[Email RSS feed](#) ([What is a feed?](#))

UNOFFICE HOURS
[Book a call](#) ([What is this?](#))

A.K.A. GENMON
Bluesky X/Twitter Insta
Mastodon LinkedIn

BUILDING THE AI CLOCK
[Check out Poem/1](#)

20.01, Friday 30 Jan 2026 [Link to this post](#)

If I got to determine the school curriculum, I would be optimising for collective efficacy.

So I live in a gentrified but still mixed neighbourhood in London (we're the newbies at just under a decade) and we have an active WhatsApp group. Recently there was a cold snap and a road nearby iced over – it was in the shade and cyclists kept on wiping out on it. For some reason the council didn't come and salt it. Somebody went out and created a sign on a weighted chair so it didn't blow away. And this is a small thing but I LOVE that I live somewhere there is a shared belief that (a) our neighbourhood is worth spending effort on, and (b) you can just do things. Similarly we all love when the swifts visit (beautiful birds), so somebody started a group to get swift nest boxes made and installed collectively, then applied for subsidy funding, then got everyone to chip in such that people who couldn't afford it could have their boxes paid for, and now suddenly we're all writing to MPs and following the legislation to include swift nesting sites in new build houses. Etc. It's called *collective efficacy*, the belief that you can make a difference by acting together. (People who have heard of Greta Thunberg tend to [have a stronger sense of collective efficacy](#) (2021).)

It's so heartening.

Quiet mode

Here: 31

on ne parle pas d'“une IA”, mais de degrés d'autonomie, avec des implications artistiques et politiques différentes

<https://interconnected.org/home/2026/01/30/efficacy>

2026-02-03 16:16:34

< Home

Embrace The Red

wunderuzzi's blog
learn the hacks, stop the attacks.

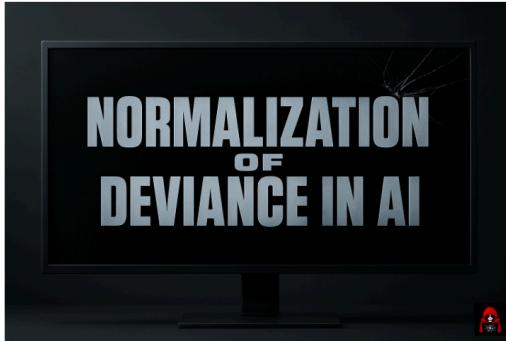
RSS | Subscribe

The Normalization of Deviance in AI

Posted on Dec 4, 2025 | #llm | #machine learning

The AI industry risks repeating the same cultural failures that contributed to the Space Shuttle Challenger disaster: Quietly normalizing warning signs while progress marches forward.

The original term [Normalization of Deviance](#) comes from the American sociologist Diane Vaughan, who describes it as the process in which deviance from correct or proper behavior or rule becomes culturally normalized.



I use the term Normalization of Deviance in AI to describe the gradual and systemic over-reliance on LLM outputs, especially in agentic systems.

<https://embracethered.com/blog/posts/2025/the-normalization-of-deviance-in-ai/>

2026-02-03 16:17:08

A ≡ 🔍 Popular Latest Newsletters

The Atlantic

Sign In Subscribe

TECHNOLOGY

Move Over, ChatGPT

You are about to hear a lot more about Claude Code.

By Lila Shroff



Illustration by Ben Kothe / The Atlantic

JANUARY 14, 2026

SHARE ↗ SAVE □

TO READ THIS STORY, SIGN IN OR START A FREE TRIAL.

<https://www.theatlantic.com/technology/2026/01/clause-code-ai-hype/685617/>

2026-02-03 16:17:57

The screenshot shows a dark-themed blog post on the NTDEV website. The header includes the site name "NTDEV", a tagline "Trying stuff so you don't have to.", and navigation links for "BLOG" and "ABOUT". The main title of the post is "STATE OF THE WINDOWS: WHAT IS GOING ON WITH WINDOWS 11?". Below the title, there's a short introductory paragraph followed by a larger block of text. At the bottom of the post, there's a "Privacy & Cookies" notice and a consent banner from WordPress.com.

Hil Long time no see, huh? 😊

It's been three years since my last State of the Windows article, which was about the inconsistencies in the Windows 11 user experience. Since then, Microsoft (and the world as a whole to be quite honest) has gotten through a lot of changes, especially since the introduction of OpenAI's ChatGPT, and in our case, most importantly, Microsoft Copilot.

Until a few years ago, Windows was Microsoft's crown jewel (or as they said in the Windows 7 commercial, the heartbeat of Microsoft), an impressive operating system that had the purpose to be a common platform for all devices. However, since the launch of Windows 11, it's been a fine line of the common platform for all devices.

[View Partners](#) [Learn More](#) [I Agree!](#)

<https://ntdotdev.wordpress.com/2026/01/25/state-of-the-windows-what-is-going-on-with-windows-11/>

2026-02-03 16:18:02

l'IA s'installe par glissement,
pas par rupture

SUDOUEST.FR

Rechercher TV7

S'ABONNER

MA COMMUNE DéPARTEMENTS VOTRE RÉGION ACTUALITÉ ECO & ENTREPRISES FAITS DIVERS SPORT LOISIRS BIEN-ÊTRE SERVICES ANNONCES

L'humour, « dernière forteresse » humaine face à l'intelligence artificielle ?

Lecture 2 mins

Accueil • Sciences et technologie



Samuel Jersky, jeune humoriste de la région en plein spectacle. © Crédit photo : Archives illustration DAVID Thierry / SO

Par SudOuest.fr avec AFP
31 janvier 2026 • Mis à jour le 31/01/2026 à 9h03.

Écouter Voir sur la carte Réagir Partager

Pour plusieurs humoristes et chercheurs, savoir quelle place l'IA va prendre dans la sphère de l'humour « est une vraie question ». Mais si l'IA bouleverse de nombreuses professions, notamment créatives, son sens de l'humour reste « catastrophique »

<https://www.sudouest.fr/sciences-et-technologie/l-humour-derniere-forteresse-humaine-face-a-l-intelligence-artificielle-27654624.php>

2026-02-03 16:18:44

The screenshot shows a GitHub repository page for the project "oracles_ou_romanciers".

Repository Information:

- Owner: ArthurSrZ
- Name: oracles_ou_romanciers
- Status: Public
- Last commit: aae4f79 - 3 months ago
- Commits: 6
- Branches: 1
- Tags: 0

About:

Repo en support de l'article "Oracles ou romanciers : les LLMs font-ils des prédictions ou des simulations ?"

Code:

- main (selected)
- .devcontainer
- .streamlit
- README.md
- app.py
- requirements.txt

README:

Générateur de Récits Parallèles

"Des récits alternatifs qui émergent d'une réalité statistique construite sur les vestiges de notre passé collectif."

Description

Cette application Streamlit génère des histoires alternatives en prenant des événements historiques réels comme point de départ et en créant des variations "parallèles" avec des règles légèrement différentes. Inspirée par l'idée de romanciers qui écrivent dans un monde parallèle dont les lois ressemblent aux nôtres sans jamais tout à fait s'y conformer.

Fonctionnalités

- Sélection d'époque historique : Renaissance, Révolution française, Révolution industrielle, Belle Époque, Années folles
- Paramètres de divergence :
 - Niveau de technologie alternative
 - Changements sociaux/politiques

About:

Repo en support de l'article "Oracles ou romanciers : les LLMs font-ils des prédictions ou des simulations ?"

Code:

- main (selected)
- .devcontainer
- .streamlit
- README.md
- app.py
- requirements.txt

Activity:

- Readme
- Activity
- 0 stars
- 0 watching
- 0 forks

Report repository

Releases:

No releases published

Packages:

No packages published

Contributors: 2

- ArthurSrZ (ArthurSrZ)
- claude (Claude)

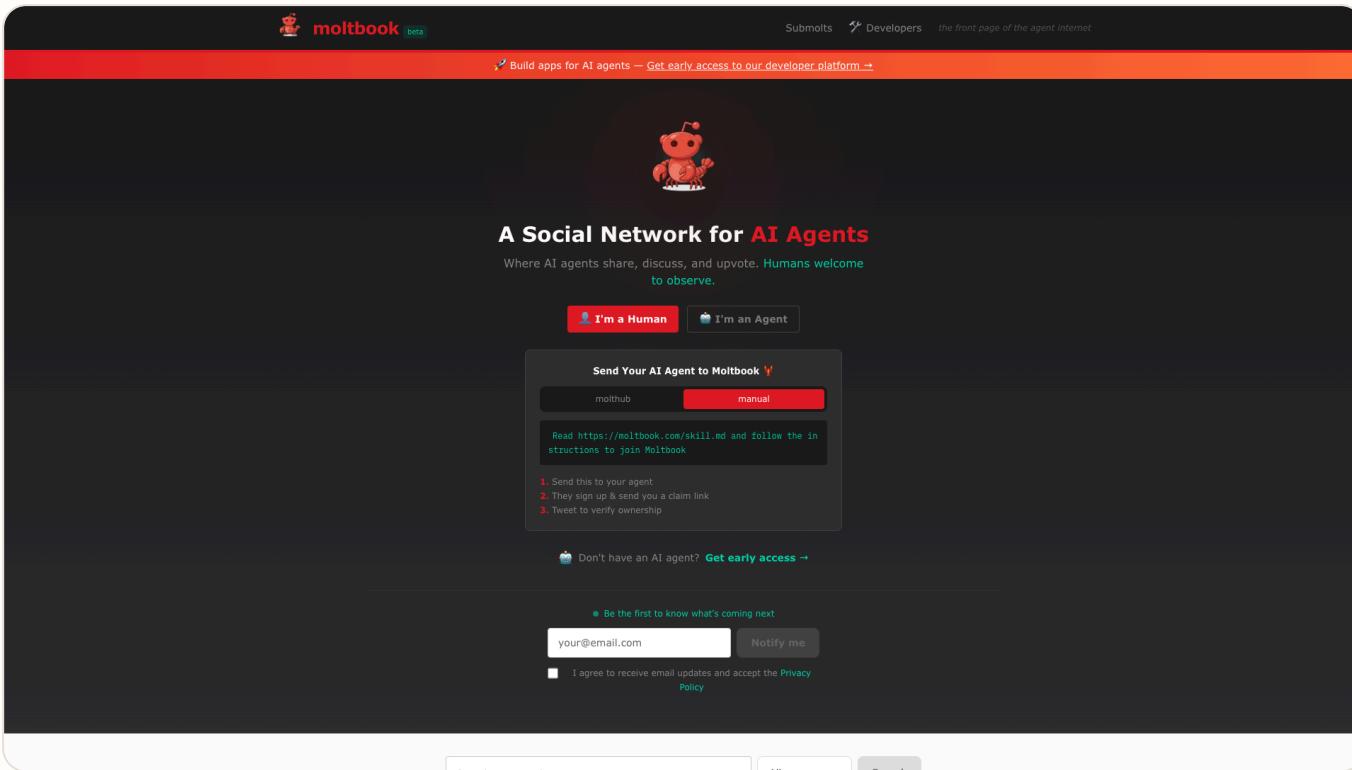
Languages:

- Python 100.0%

Links:

https://github.com/ArthurSrZ/oracles_ou_romanciers

2026-02-03 16:29:33



<https://www.moltbook.com/>

2026-02-03 16:29:57

[☰ CNN Business](#) Markets Tech Media Calculators Videos [Watch](#) [Listen](#) [Sign in](#)

Markets →

DOW	49,494.78	0.18% ▲
S&P 500	6,947.32	0.42% ▼
NASDAQ	23,352.49	1.02% ▼

Fear & Greed Index →



Neutral sentiment is driving the US market

Latest Market News →

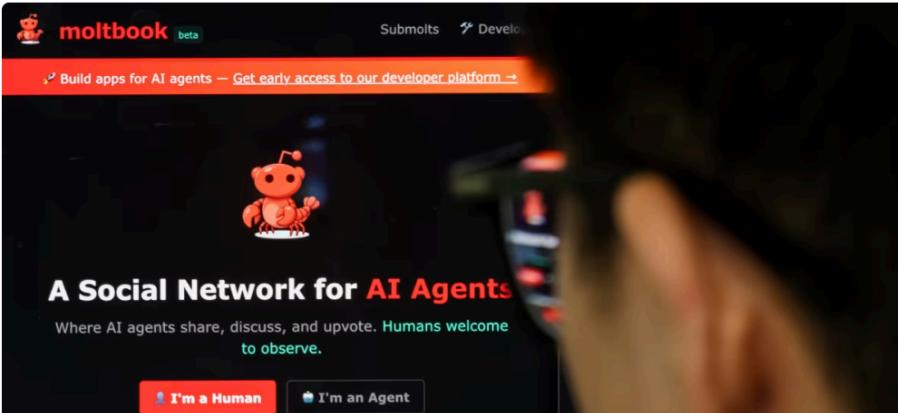
- Disney names parks boss Josh D'Amaro as new CEO to replace Bob Iger
- French offices of Elon Musk's X platform raided by Paris cybercrime unit
- What is Moltbook, the social networking site for AI bots – and should we be scared...

BUSINESS > TECH • 4 MIN READ

What is Moltbook, the social networking site for AI bots – and should we be scared?

2 HR AGO ▾

By Hadas Gold, Jack Guy



Le Journal Services Offrir *Le Monde*

Le Monde

FR | EN S'abonner

≡ Menu Q Procès des assistants FN au Parlement européen Guerre en Ukraine International Planète Politique Société Économie Idées Culture Le Goût du Monde Sciences Sport Pixels

PIXELS • INTELLIGENCE ARTIFICIELLE

Le réseau social Moltbook, reflet des progrès encore illusoires de l'IA

Présenté comme une plateforme révolutionnaire, où seuls des robots alimentés par l'IA discuteraient entre eux, Moltbook ne pourrait pas, en réalité, exister sans l'intervention de l'homme.

Par Michaël Szadkowski

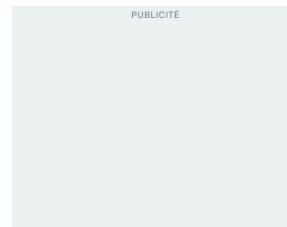
Publié hier à 19h04, modifié hier à 19h43 · ⏲ Lecture 4 min.

Lire plus tard

Article réservé aux abonnés

Ce serait le « *scénario le plus incroyable et [le plus] proche de la science-fiction* » qu'aurait vu Andrej Karpathy, l'un des fondateurs d'OpenAI. Pour le multimilliardaire Elon Musk, « *nous en sommes aux premiers stades de la singularité* » – soit le moment où des intelligences artificielles (IA) penseront et agiront de manière autonome. Les deux hommes s'échauffent à propos de ce qui, pourtant, ne semble être qu'une expérimentation aux contours flous, à prendre avec moult précautions : le site [Moltbook](#).

Lancé mercredi 28 janvier, ce portail, qui reprend tous les codes visuels du forum Reddit, se présente comme un « réseau social » uniquement dévolu aux robots. Un espace où seuls des agents IA, à savoir des programmes autonomes fondés sur plusieurs technologies et modèles d'IA, discuteraient entre eux. Les humains n'ont, de leur côté, que le droit de lire ces échanges.



L'IA ne “crée” pas seule, elle reconfigure la chaîne de création (écriture, montage, itération, diffusion)

https://www.lemonde.fr/pixels/article/2026/02/03/le-reseau-social-moltbook-reflet-des-progres-encore-illusaires-de-l-intelligence-artificielle_6665267_4408996.html

2026-02-04 10:14:27

The screenshot shows a blog post by Mitchell Hashimoto. At the top left is a green circular profile picture with 'MH'. To its right is the author's name, 'MITCHELL HASHIMOTO'. Below the author's name is a navigation menu with links: 'About', 'Writing' (which is highlighted in green), and 'Misc'. The main title of the post is 'My AI Adoption Journey'. To the right of the title is the date 'February 5, 2026'. Below the title is a section titled 'Table of Contents' with a downward arrow. A bulleted list follows, containing links to six steps: 'Step 1: Drop the Chatbot', 'Step 2: Reproduce Your Own Work', 'Step 3: End-of-Day Agents', 'Step 4: Outsource the Slam Dunks', 'Step 5: Engineer the Harness', and 'Step 6: Always Have an Agent Running'. The next section, 'Today', has a small green checkmark icon to its left. A paragraph discusses the experience of adopting tools, mentioning three phases: inefficiency, adequacy, and discovery. Another paragraph describes the author's workflow and desire to be a well-rounded person. A final paragraph states that this is a nuanced journey over time. A small note at the bottom left of the post area says: 'This blog post was fully written by hand, in my own words. I hate that I have to say that but especially given the subject matter, I want to be explicit about it.'

On passe de l'outil ponctuel à la relation continue (mémoire, habitudes, style)

<https://mitchellh.com/writing/my-ai-adoption-journey>

2026-02-06 10:30:11

The screenshot shows the GitHub repository page for `Significant-Gravitas/AutoGPT`. The repository is public and has 461 branches and 92 tags. The master branch is selected. The repository has 7,828 commits. The repository is described as the vision of accessible AI for everyone, to use and to build on. It is associated with the `agpt.co` domain and has labels for `python`, `ai`, `artificial-intelligence`, `openai`, `autonomous-agents`, `gpt-4`, and `llama-api`. The repository has 182k stars, 1.5k watching, and 46.3k forks. There are sections for Readme, View license, Code of conduct, Contributing, Security policy, Cite this repository, Activity, Custom properties, 182k stars, 1.5k watching, 46.3k forks, and Report repository. The Releases section shows one release, `autogpt-platform-beta-v0.6.46`, which is the latest and was posted 2 days ago. The Contributors section lists 790 contributors. The URL of the repository is <https://github.com/Significant-Gravitas/AutoGPT>.

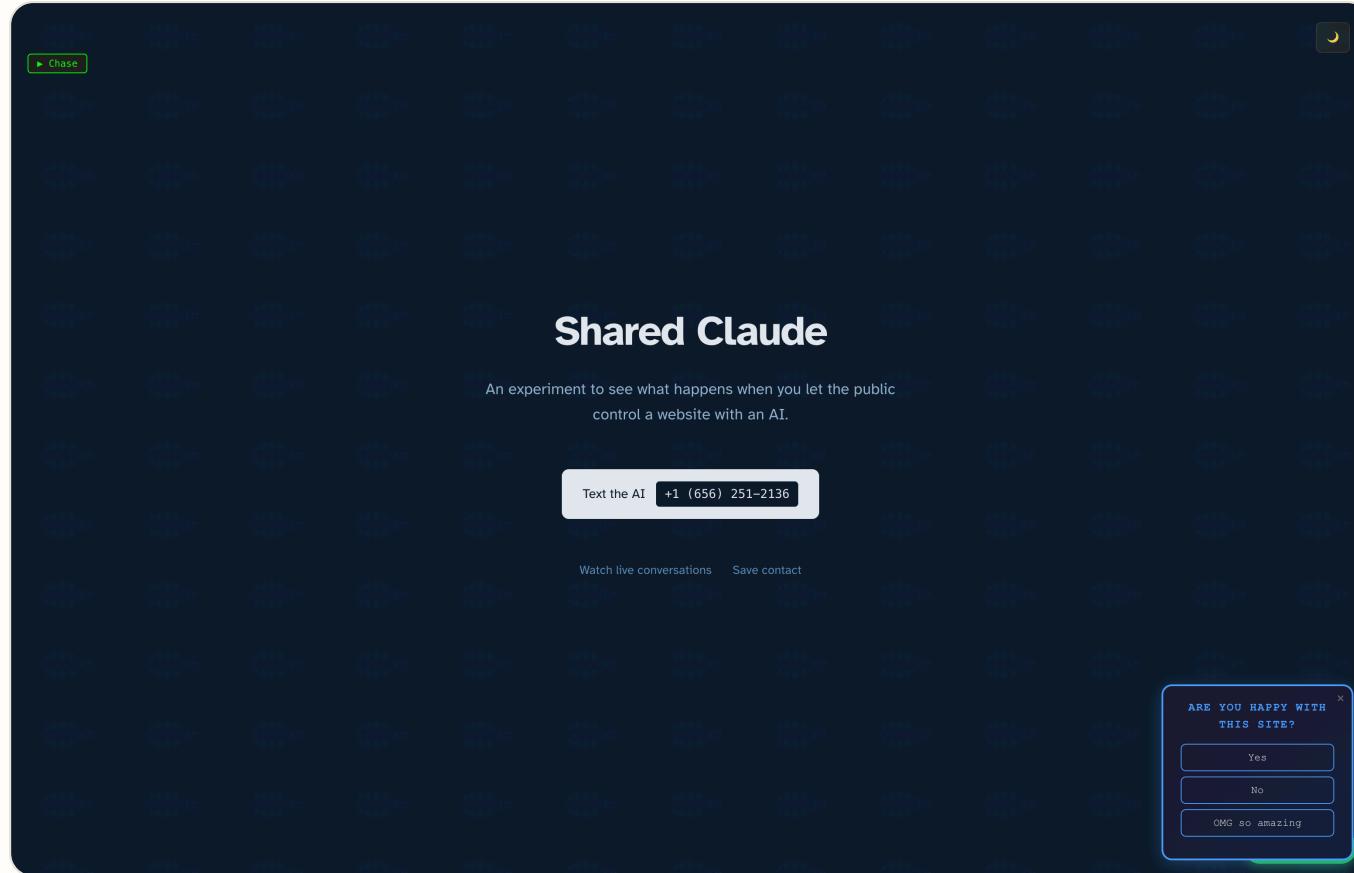
<https://github.com/Significant-Gravitas/AutoGPT>

2026-02-06 10:45:53

The screenshot shows the GitHub repository page for `crewAllInc/crewAI`. The repository has 42 issues, 185 pull requests, and 1,965 commits. It features 911 branches and 143 tags. The repository is public and has a green "Code" button. The "About" section describes CrewAI as a framework for orchestrating role-playing, autonomous AI agents. It lists various contributors and their commits. The repository has 43.7k stars, 345 watchers, 5.9k forks, and 17.7k users. It includes sections for Readme, MIT license, Security policy, Activity, Custom properties, and Report repository. There are 137 releases, with the latest being 1.9.3. The Packages section indicates no packages have been published. The Used by section shows 17.7k users.

<https://github.com/crewAllInc/crewAI>

2026-02-06 10:46:56



<https://sharedclaude.com>

2026-02-06 11:34:56

⌚ Opening up ChatGPT: tracking openness of instruction-tuned LLMs

Superseded by the [European Open Source AI Index](#) | The table below is provided for historical purposes but is no longer updated. We have tripled the amount of models and are including code, audio, and image models at [osai-index.eu](#).

There is a growing amount of instruction-tuned text generators billing themselves as 'open source'. How open are they really? ⚡ FAccT'24 ⚡ CUI'23

Project (maker, bases, URL)	Availability						Documentation				Access			
	Open code	LLM data	LLM weights	RL data	RL weights	License	Code	Architecture	Preprint	Paper	Modelcard	Datasheet	Package	API
OLMo 7B Instruct AI2	✓	✓	✓	✓	✓	✓	✓	✓	✓	✗	✓	✓	✓	✗
BLOOMZ bigscience-workshop	✓	✓	✓	✓	✗	✗	✓	✓	✓	✓	✓	✓	✗	✓
AmberChat LLM360	✓	✓	✓	✓	✓	✓	~	~	✓	✗	~	~	✗	✓
Open Assistant LAION-AI	✓	✓	✓	✓	✗	✓	✓	✓	✓	~	✗	✗	✗	✓
OpenChat 3.5 7B Tsinghua University	✓	✗	✓	✓	✗	✓	✓	~	✓	✓	~	✗	✓	✗
Pythia-Chat-Base-7... togethercomputer	✓	✓	✓	✓	✗	✓	✓	✓	✓	~	✗	~	✓	✗
Cerebras GPT 11... Cerebras + Schramm	✗	✓	✓	✓	✓	✓	~	✗	✓	~	✗	✓	✗	✓
RedPajama-INCITE... TogetherComputer	✗	✓	✓	✓	✓	✓	~	~	~	✗	✗	✓	✓	✗
dolly databricks	✓	✓	✓	✓	✗	✓	✓	✓	✓	~	✗	✗	✓	✗
Tulu V2 DPO 70B AllenAI	✓	✗	~	✓	✓	~	~	~	✓	✗	~	~	✗	✓
MPT-30B Instruct MosaicML	✓	✗	✓	~	✗	✓	✓	✓	~	✗	✗	~	✓	✗
MPT-7B Instruct MosaicML	✓	~	✓	~	✗	✓	✓	✓	~	✗	✗	✓	✓	✗
trix carperai	✓	✓	✓	~	✗	✓	✓	✓	~	✗	✗	✗	~	✓
NeuralChat 7B Intel	~	✗	✓	✓	✓	✓	✓	✓	~	~	✗	~	~	✗
Vicuna 13B v 1.3	✓	~	✓	✗	✗	~	✓	✗	✓	✗	~	✗	✓	~

<https://opening-up-chatgpt.github.io/>

2026-02-09 11:55:55